Chen Hongyu

chenhongyu2048@sjtu.edu.cn • chenhongyu2048@gmail.com • (+86) 156-9529-5531
Citizenship: China

Research interests

System for Machine Learning, i.e., MLSys

Education

2024 – Present **Shanghai Jiao Tong University** – Minhang, Shanghai
Ph.D. student in Computer Science and Technology
School of Computer Science
Advisor: Professor Shixuan Sun.

2021 – 2024 **National University of Defense Technology** – Changsha, Hunan
M.S. in Computer Technology
School of Computer Science
Advisor: Professor Dongsheng Li.

2017 – 2021 **Nanjing University** – Nanjing, Jiangsu
B.S. in Space Science and Technology
School of Astronomy and Space Science.

Honors and scholarships

2020 The Renmin Schoolarship (Nanjing University)

2019 1st Prize (The 11st National College Studenst Mathematical Competition)

2019 3rd Prize (The 22nd Forum of Sciences and Arts at Nanjing University)

2016 2nd Prize Provincial, Chinese High School Mathematics League

Publications

Published:

Nov. 2026 **InfiniLoRA: Disaggregated Multi-LoRA Serving for Large Language Models**
Hongyu Chen, Letian Ruan, Zilin Xu, Yuchen Li, Xinyu Chen, Bingsheng He, Jingwen Leng, Jieru Zhao, Shixuan Sun, Minyi Guo.
We introduce a shared LoRA Server to provide efficient Multi-LoRA service, equipped with parallelism-aware execution, SLO-driven provisioning, and critical-path optimizations, achieving an average 3.05× increase in serviceable request rate.
The 2026 ACM SIGOPS Annual Technical Conference (ATC, CCF-A).
Camera-ready: *arXiv:2604.07173v2*

Aug. 2023 **Auto-Divide GNN: Accelerating GNN Training with Subgraph Division**
Hongyu Chen, Zhejiang Ran, Keshi Ge, Zhiquan Lai, Jingfei Jiang, Dongsheng Li.
The 29th International European Conference on Parallel and Distributed Computing (Euro-Par, CCF-B).

Preprints:

Nov. 2026 **LiquidCP: Decoupling Attention Tasks from Communication and Kernel Boundaries in Context Parallelism**
Hongyu Chen, Yikang Ruan, Yuhai Long, Ziyu Huang, Yiyu Luo, Liukun Yu, Yangjie Zhou, Zihan Liu, Weihao Cui, Zijun Li, Anbang Wu, Yu Feng, Jing Yang, Jingwen Leng, Bingsheng He, Minyi Guo, Shixuan Sun.
On eight H100 GPUs, MegaRing achieves 1.23× forward and 1.18× backward core-attention speedups; MegaDCP achieves 1.20× for decode-only batches and 1.35× for mixed batches.
arXiv:2611.04401

Experience

Jun. 2026 – Present **Research Intern at Huawei** – Shanghai
AI Basic Algorithm Research Department, Huawei Consumer Cloud.
Research and development of efficient training systems for multimodal large language models (MLLMs).

Nov. 2025 – Jan. 2026 **Research Intern at ByteDance** – Shanghai
Seed-Infra, LLM Inference Team.
AI infrastructure research and development, focusing on efficient Multi-LoRA serving for large language models.

Jul. 2023 – Oct. 2023 **Research Intern at Huawei** – Hangzhou
Advanced Storage and Computing Lab, Central Research Institute.
C++ software development for Huawei's graph computing engine HyG, responsible for the development of graph neural network operators.

2022 Autumn **Teaching Assistant** – National University of Defense Technology
Course: *Data Structures and Algorithms.*
Responsible for answering students' questions and preparing examinations.